

Evaluating Student Writing

We aim to identify elements in student writing i.e. we segment text and classify argumentative and rhetorical elements i.e. predict human annotations in essays written by 6th-12th grade students.

(The demo and visualisation on the test data is at the end of the file)

Setup

```
In [1]: # Install libraries to get evaluation metrics for training data
!pip install segeval
!pip install segeval -qq
!pip install wandb
!pip install --upgrade wandb -qq
import wandb
import warnings
warnings.filterwarnings('ignore')
warnings.simplefilter('ignore')

# visualization with displacy
import pandas as pd
import os
from pathlib import Path
import spacy
from spacy import displacy
from pylab import cm, matplotlib
```

Collecting segeval

Downloading segeval-1.2.2.tar.gz (43 kB)

|██| 43 kB 178 kB/s

Preparing metadata (setup.py) ... done

Requirement already satisfied: numpy>=1.14.0 in /opt/conda/lib/python3.7/site-packages (from segeval) (1.19.5)

Requirement already satisfied: scikit-learn>=0.21.3 in /opt/conda/lib/python3.7/site-packages (from segeval) (0.23.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.21.3->segeval) (3.0.0)

Requirement already satisfied: scipy>=0.19.1 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.21.3->segeval) (1.7.2)

Requirement already satisfied: joblib>=0.11 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.21.3->segeval) (1.1.0)

Building wheels for collected packages: segeval

Building wheel for segeval (setup.py) ... done

Created wheel for segeval: filename=segeval-1.2.2-py3-none-any.whl size=16181 sha256=5c2e0aa8849cf0347e9ddf29b82f81ec1962982dbec2d2ec7c8ee4e52d34e745

Stored in directory: /root/.cache/pip/wheels/05/96/ee/7cac4e74f3b19e3158dce26a20alc86b3533c43ec72a549fd7

Successfully built segeval

Installing collected packages: segeval

Successfully installed segeval-1.2.2

WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: <https://pip.pypa.io/warnings/venv>

Requirement already satisfied: wandb in /opt/conda/lib/python3.7/site-packages (0.12.7)

Requirement already satisfied: promise<3,>=2.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (2.3)

Requirement already satisfied: yaspin>=1.0.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (2.1.0)

Requirement already satisfied: sentry-sdk>=1.0.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (1.5.0)

Requirement already satisfied: PyYAML in /opt/conda/lib/python3.7/site-packages (from wandb) (6.0)

Requirement already satisfied: Click!=8.0.0,>=7.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (8.0.3)

Requirement already satisfied: psutil>=5.0.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (5.8.0)

Requirement already satisfied: six>=1.13.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (1.16.0)

Requirement already satisfied: protobuf>=3.12.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (3.19.1)

Requirement already satisfied: requests<3,>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (2.25.1)

Requirement already satisfied: GitPython>=1.0.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (3.1.24)

Requirement already satisfied: docker-pycreds>=0.4.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (0.4.0)

Requirement already satisfied: python-dateutil>=2.6.1 in /opt/conda/lib/python3.7/site-packages (from wandb) (2.8.0)

Requirement already satisfied: subprocess32>=3.5.3 in /opt/conda/lib/python3.7/site-packages (from wandb) (3.5.4)

Requirement already satisfied: configparser>=3.8.1 in /opt/conda/lib/python3.7/site-packages (from wandb) (5.1.0)

Requirement already satisfied: pathtools in /opt/conda/lib/python3.7/site-packages (from wandb) (0.1.2)

Requirement already satisfied: shortuuid>=0.5.0 in /opt/conda/lib/python3.7/site-packages (from wandb) (1.0.8)

Requirement already satisfied: importlib-metadata in /opt/conda/lib/python3.7/site-packages (from Click!=8.0.0,>=7.0->wandb) (4.8.2)

Requirement already satisfied: gitdb<5,>=4.0.1 in /opt/conda/lib/python3.7/site-packages (from GitPython>=1.0.0->wandb) (4.0.9)

Requirement already satisfied: typing-extensions>=3.7.4.3 in /opt/conda/lib/python3.7/site-packages (from GitPython>=1.0.0->wandb) (3.10.0.2)

Requirement already satisfied: chardet<5,>=3.0.2 in /opt/conda/lib/python3.7/site-packages (from requests<3,>=2.0.0->wandb) (4.0.0)

Requirement already satisfied: idna<3,>=2.5 in /opt/conda/lib/python3.7/site-packages (from requests<3,>=2.0.0->wandb) (2.10)

Requirement already satisfied: urllib3<1.27,>=1.21.1 in /opt/conda/lib/python3.7/site-packages (from requests<3,>=2.0.0->wandb) (1.26.7)

Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.7/site-packages (from requests<3,>=2.0.0->wandb) (2021.10.8)

Requirement already satisfied: termcolor<2.0.0,>=1.1.0 in /opt/conda/lib/python3.7/site-packages (from yaspin>=1.0.0->wandb) (1.1.0)

Requirement already satisfied: smmap<6,>=3.0.1 in /opt/conda/lib/python3.7/site-packages (from gitdb<5,>=4.0.1->GitPython>=1.0.0->wandb) (3.0.5)

Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.7/site-packages (from importlib-metadata->Click!=8.0.0,>=7.0->wandb) (3.6.0)

WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: <https://pip.pypa.io/warnings/venv>

Choose Model to run

```
In [2]: print("Choose the model to run: \n 1. LongFormer \n 2. BigBird")
print("Kindly enter 1 or 2")
while True:
    input_model = input()
    if len(input_model)>1:
        print("Invalid Input, Kindly enter 1 or 2")
    else:
        if input_model == '1':
            model_checkpoint = "allenai/longformer-base-4096"
            print("The model choosen is Longformer")
            break
        elif input_model == '2':
            model_checkpoint = "google/bigbird-roberta-base"
            print("The model choosen is BigBird")
            break
        else:
            print("Invalid Input, Kindly enter 1 or 2")
```

Choose the model to run:

1. LongFormer
2. BigBird

Kindly enter 1 or 2

The model choosen is Longformer

Configurations

```
In [3]: # Configurations

SAMPLE = False # Used for debugging

EXP_NUM = 4
```

```

task = "ner"
max_length = 1024
stride = 128
min_tokens = 6
model_path = f'{model_checkpoint.split("/")[-1]}-{EXP_NUM}'
max_length = 1024
batch_size = 4

# TRAINING HYPERPARAMS
BS = 4
GRAD_ACC = 8
LR = 5e-5
WD = 0.01
WARMUP = 0.1
N_EPOCHS = 5

```

Data Preprocessing

```

In [4]: import pandas as pd

# Importing the train data
train = pd.read_csv('../input/feedback-prize-2021/train.csv')
train.head()

```

```

Out[4]:

```

	id	discourse_id	discourse_start	discourse_end	discourse_text	discourse_type
0	423A1CA112E2	1.622628e+12	8.0	229.0	Modern humans today are always on their phone....	Leac
1	423A1CA112E2	1.622628e+12	230.0	312.0	They are some really bad consequences when stu...	Positior
2	423A1CA112E2	1.622628e+12	313.0	401.0	Some certain areas in the United States ban ph...	Evidence
3	423A1CA112E2	1.622628e+12	402.0	758.0	When people have phones, they know about certa...	Evidence
4	423A1CA112E2	1.622628e+12	759.0	886.0	Driving is one of the way how to get around. P...	Claim

```

In [5]: # Viewing the unique classes in the dataset
classes = train.discourse_type.unique().tolist()
classes

```

```
Out[5]: ['Lead',
        'Position',
        'Evidence',
        'Claim',
        'Concluding Statement',
        'Counterclaim',
        'Rebuttal']
```

```
In [6]: # Setting label incides
from collections import defaultdict
tags = defaultdict()

for i, c in enumerate(classes):
    print(i,c)
    tags[f'B-{c}'] = i
    tags[f'I-{c}'] = i + len(classes)
tags[f'O'] = len(classes) * 2
tags[f'Special'] = -100

l2i = dict(tags)

i2l = defaultdict()
for k, v in l2i.items():
    i2l[v] = k
i2l[-100] = 'Special'

i2l = dict(i2l)

N_LABELS = len(i2l) - 1
```

```
0 Lead
1 Position
2 Evidence
3 Claim
4 Concluding Statement
5 Counterclaim
6 Rebuttal
```

```
In [7]: # Viewing the tags assigned to the classes
tags
```

```
Out[7]: defaultdict(None,
                    {'B-Lead': 0,
                     'I-Lead': 7,
                     'B-Position': 1,
                     'I-Position': 8,
                     'B-Evidence': 2,
                     'I-Evidence': 9,
                     'B-Claim': 3,
                     'I-Claim': 10,
                     'B-Concluding Statement': 4,
                     'I-Concluding Statement': 11,
                     'B-Counterclaim': 5,
                     'I-Counterclaim': 12,
                     'B-Rebuttal': 6,
                     'I-Rebuttal': 13,
                     'O': 14,
                     'Special': -100})
```

```
In [8]: # Reading raw text from the essay files
```

```
from pathlib import Path

path = Path('../input/feedback-prize-2021/train')

def get_raw_text(ids):
    with open(path/f'{ids}.txt', 'r') as file: data = file.read()
    return data
```

```
In [9]: # Grouping annotations based on discourse_type, discourse_start, discourse_end
# predictionstring to form single tuple for each essay

df1 = train.groupby('id')['discourse_type'].apply(list).reset_index(name='class')
df2 = train.groupby('id')['discourse_start'].apply(list).reset_index(name='start')
df3 = train.groupby('id')['discourse_end'].apply(list).reset_index(name='ends')
df4 = train.groupby('id')['predictionstring'].apply(list).reset_index(name='predictionstring')

#Merging the dataframes
df = pd.merge(df1, df2, how='inner', on='id')
df = pd.merge(df, df3, how='inner', on='id')
df = pd.merge(df, df4, how='inner', on='id')

#Adding raw essay text to the merged data
df['text'] = df['id'].apply(get_raw_text)

In [10]: #Viewing the merged data
df.head()
```

Out[10]:

	id	classlist	starts	ends	predictionstrings	text
0	0000D23A521A	[Position, Evidence, Evidence, Claim, Countercl...	[0.0, 170.0, 358.0, 438.0, 627.0, 722.0, 836.0...	[170.0, 357.0, 438.0, 626.0, 722.0, 836.0...	[0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 1...	Some people believe that the so called "face" o...
1	00066EA9880D	[Lead, Position, Claim, Evidence, Claim, Evide...	[0.0, 456.0, 638.0, 738.0, 1399.0, 1488.0, 231...	[455.0, 592.0, 738.0, 1398.0, 1487.0, 2219.0, ...	[0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 1...	Driverless cars are exactly what you would exp...
2	000E6DE9E817	[Position, Counterclaim, Rebuttal, Evidence, C...	[17.0, 64.0, 158.0, 310.0, 438.0, 551.0, 776.0...	[56.0, 157.0, 309.0, 422.0, 551.0, 775.0, 961....	[2 3 4 5 6 7 8, 10 11 12 13 14 15 16 17 18 19 ...	Dear: Principal\n\nI am arguing against the po...
3	001552828BD0	[Lead, Evidence, Claim, Claim, Evidence, Claim...	[0.0, 161.0, 872.0, 958.0, 1191.0, 1542.0, 161...	[160.0, 872.0, 957.0, 1190.0, 1541.0, 1612.0, ...	[0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 1...	Would you be able to give your car up? Having ...
4	0016926B079C	[Position, Claim, Claim, Claim, Claim, Evidenc...	[0.0, 58.0, 94.0, 206.0, 236.0, 272.0, 542.0, ...	[57.0, 91.0, 150.0, 235.0, 271.0, 542.0, 650.0...	[0 1 2 3 4 5 6 7 8 9, 10 11 12 13 14 15, 16 17...	I think that students would benefit from learn...

In [11]: *# Size of the dataset before removing outliers*

```
df2=df
df.shape
```

Out[11]: (15594, 6)

In [12]: *# Removing outliers i.e. essays with more than 5 occurrences of the same class*

```
from collections import Counter
res = []
for i in range(len(df['classlist'])):
    temp = df['classlist'][i]
    res.append(dict(Counter(temp)))

def countOccurrences(res, df2):
    df = pd.DataFrame(res)
    classes = ['Lead', 'Position', 'Evidence', 'Claim', 'Concluding Statement',
    for c in classes:
        index = df[df[c]>5].index
```

```

df = df.drop(index)
df2 = df2.drop(index)

return df2

df2 = countOccurrences(res,df2)
df = df2

```

```

In [13]: #Size of the dataset after removing outliers
df.shape

```

```

Out[13]: (12736, 6)

```

```

In [14]: # debugging
if SAMPLE: df = df.sample(n=10).reset_index(drop=True)

```

```

In [15]: # Performing Train Test split
from datasets import Dataset, load_metric

ds = Dataset.from_pandas(df)
datasets = ds.train_test_split(test_size=0.1, shuffle=True, seed=42)
datasets

```

```

Out[15]: DatasetDict({
  train: Dataset({
    features: ['id', 'classlist', 'starts', 'ends', 'predictionstrings',
'text', '__index_level_0__'],
    num_rows: 11462
  })
  test: Dataset({
    features: ['id', 'classlist', 'starts', 'ends', 'predictionstrings',
'text', '__index_level_0__'],
    num_rows: 1274
  })
})

```

```

In [16]: #Initialing the tokenizer
from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained(model_checkpoint, add_prefix_space=True)

```

```

Downloading: 0%|          | 0.00/694 [00:00<?, ?B/s]
Downloading: 0%|          | 0.00/878k [00:00<?, ?B/s]
Downloading: 0%|          | 0.00/446k [00:00<?, ?B/s]
Downloading: 0%|          | 0.00/1.29M [00:00<?, ?B/s]

```

```

In [17]: # If a span is created without a starting token for a class
# then we convert the first token to be the starting token
def fix_beginnings(labels):
    for i in range(1,len(labels)):
        curr_lab = labels[i]
        prev_lab = labels[i-1]
        if curr_lab in range(7,14):
            if prev_lab != curr_lab and prev_lab != curr_lab - 7:
                labels[i] = curr_lab - 7
    return labels

```

```

In [18]: # tokenizing and adding labels
def tokenize_and_align_labels(examples):

```



```

o = tokenizer(examples['text'], truncation=True, padding=True,
              return_offsets_mapping=True, max_length=max_length,
              stride=stride, return_overflowing_tokens=True)
#print(o.keys())
sample_mapping = o["overflow_to_sample_mapping"]
offset_mapping = o["offset_mapping"]

o["labels"] = []

for i in range(len(offset_mapping)):

    sample_index = sample_mapping[i]

    labels = [l2i['O'] for i in range(len(o['input_ids'][i]))]

    for label_start, label_end, label in \
    list(zip(examples['starts'][sample_index], examples['ends'][sample_index],
            examples['labels'][sample_index])):
        for j in range(len(labels)):
            token_start = offset_mapping[i][j][0]
            token_end = offset_mapping[i][j][1]
            if token_start == label_start:
                labels[j] = l2i[f'B-{label}']
            if token_start > label_start and token_end <= label_end:
                labels[j] = l2i[f'I-{label}']

    for k, input_id in enumerate(o['input_ids'][i]):
        if input_id in [0,1,2]:
            labels[k] = -100

    labels = fix_beginnings(labels)
    o["labels"].append(labels)

return o

```

```

In [19]: # Tokenising both train and test
tokenized_datasets = datasets.map(tokenize_and_align_labels, batched=True, \
                                  batch_size=20000,
                                  remove_columns=datasets["train"].column_names)

0%|          | 0/1 [00:00<?, ?ba/s]
0%|          | 0/1 [00:00<?, ?ba/s]

```

```

In [20]: tokenized_datasets

```

```

Out[20]: DatasetDict({
  train: Dataset({
    features: ['attention_mask', 'input_ids', 'labels', 'offset_mapping',
'overflow_to_sample_mapping'],
    num_rows: 11778
  })
  test: Dataset({
    features: ['attention_mask', 'input_ids', 'labels', 'offset_mapping',
'overflow_to_sample_mapping'],
    num_rows: 1313
  })
})

```

Model and Training

```
In [21]: from transformers import AutoModelForTokenClassification, TrainingArguments, Tr
```

```
model = AutoModelForTokenClassification.from_pretrained(model_checkpoint, num_l
```

Downloading: 0%| | 0.00/570M [00:00<?, ?B/s]

Some weights of the model checkpoint at allenai/longformer-base-4096 were not used when initializing LongformerForTokenClassification: ['lm_head.layer_norm.bias', 'lm_head.dense.weight', 'lm_head.bias', 'lm_head.layer_norm.weight', 'lm_head.dense.bias', 'lm_head.decoder.weight']

- This IS expected if you are initializing LongformerForTokenClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing LongformerForTokenClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of LongformerForTokenClassification were not initialized from the model checkpoint at allenai/longformer-base-4096 and are newly initialized: ['classifier.bias', 'classifier.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
In [22]: #print(model)
```

```
In [23]: model_name = model_checkpoint.split("/")[-1]
args = TrainingArguments(
    f"{model_name}-finetuned-{task}",
    evaluation_strategy = "epoch",
    logging_strategy = "epoch",
    save_strategy = "epoch",
    learning_rate=LR,
    per_device_train_batch_size=BS,
    per_device_eval_batch_size=BS,
    num_train_epochs=N_EPOCHS,
    weight_decay=WD,
    report_to='wandb',
    gradient_accumulation_steps=GRAD_ACC,
    warmup_ratio=WARMUP
)
```

```
In [24]: from transformers import DataCollatorForTokenClassification

data_collator = DataCollatorForTokenClassification(tokenizer)
```

```
In [25]: # Loading Metric
metric = load_metric("sequeval")
```

Downloading: 0%| | 0.00/2.48k [00:00<?, ?B/s]

```
In [26]: import numpy as np

def compute_metrics(p):
    predictions, labels = p
    predictions = np.argmax(predictions, axis=2)

    # Remove special tokens
    true_predictions = [
```

```

[i2l[p] for (p, l) in zip(prediction, label) if l != -100]
for prediction, label in zip(predictions, labels)
]
true_labels = [
[i2l[l] for (p, l) in zip(prediction, label) if l != -100]
for prediction, label in zip(predictions, labels)
]

results = metric.compute(predictions=true_predictions, references=true_labels)
return {
    "precision": results["overall_precision"],
    "recall": results["overall_recall"],
    "f1": results["overall_f1"],
    "accuracy": results["overall_accuracy"],
}

```

```

In [27]: trainer = Trainer(
    model,
    args,
    train_dataset=tokenized_datasets["train"],
    eval_dataset=tokenized_datasets["test"],
    data_collator=data_collator,
    tokenizer=tokenizer,
    compute_metrics=compute_metrics,
)

```

```

In [28]: import os
os.environ["WANDB_DISABLED"] = "true"

```

```

In [29]: trainer.train()
wandb.finish()

```

The following columns in the training set don't have a corresponding argument in `LongformerForTokenClassification.forward` and have been ignored: offset_mapping, overflow_to_sample_mapping.

***** Running training *****

Num examples = 11778

Num Epochs = 5

Instantaneous batch size per device = 4

Total train batch size (w. parallel, distributed & accumulation) = 32

Gradient Accumulation steps = 8

Total optimization steps = 1840

Automatic Weights & Biases logging enabled, to disable set os.environ["WANDB_DISABLED"] = "true"

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

Tracking run with wandb version 0.12.16

W&B syncing is set to `offline` in this directory.

Run `wandb online` or set WANDB_MODE=online to enable cloud syncing.

[1840/1840 3:59:11, Epoch 4/5]

Epoch	Training Loss	Validation Loss	Precision	Recall	F1	Accuracy
0	1.028900	0.654350	0.102136	0.217217	0.138941	0.787744
1	0.621500	0.605735	0.143180	0.284232	0.190431	0.797515
2	0.525700	0.593650	0.159190	0.299187	0.207810	0.805811
3	0.449300	0.618323	0.173667	0.313020	0.223393	0.802356
4	0.392300	0.641289	0.171629	0.322740	0.224090	0.799704

```

The following columns in the evaluation set don't have a corresponding argument in `LongformerForTokenClassification.forward` and have been ignored: offset_mapping, overflow_to_sample_mapping.
***** Running Evaluation *****
    Num examples = 1313
    Batch size = 4
Saving model checkpoint to longformer-base-4096-finetuned-ner/checkpoint-368
Configuration saved in longformer-base-4096-finetuned-ner/checkpoint-368/config.json
Model weights saved in longformer-base-4096-finetuned-ner/checkpoint-368/pytorch_model.bin
tokenizer config file saved in longformer-base-4096-finetuned-ner/checkpoint-368/tokenizer_config.json
Special tokens file saved in longformer-base-4096-finetuned-ner/checkpoint-368/special_tokens_map.json
The following columns in the evaluation set don't have a corresponding argument in `LongformerForTokenClassification.forward` and have been ignored: offset_mapping, overflow_to_sample_mapping.
***** Running Evaluation *****
    Num examples = 1313
    Batch size = 4
Saving model checkpoint to longformer-base-4096-finetuned-ner/checkpoint-736
Configuration saved in longformer-base-4096-finetuned-ner/checkpoint-736/config.json
Model weights saved in longformer-base-4096-finetuned-ner/checkpoint-736/pytorch_model.bin
tokenizer config file saved in longformer-base-4096-finetuned-ner/checkpoint-736/tokenizer_config.json
Special tokens file saved in longformer-base-4096-finetuned-ner/checkpoint-736/special_tokens_map.json
The following columns in the evaluation set don't have a corresponding argument in `LongformerForTokenClassification.forward` and have been ignored: offset_mapping, overflow_to_sample_mapping.
***** Running Evaluation *****
    Num examples = 1313
    Batch size = 4
Saving model checkpoint to longformer-base-4096-finetuned-ner/checkpoint-1104
Configuration saved in longformer-base-4096-finetuned-ner/checkpoint-1104/config.json
Model weights saved in longformer-base-4096-finetuned-ner/checkpoint-1104/pytorch_model.bin
tokenizer config file saved in longformer-base-4096-finetuned-ner/checkpoint-1104/tokenizer_config.json
Special tokens file saved in longformer-base-4096-finetuned-ner/checkpoint-1104/special_tokens_map.json
The following columns in the evaluation set don't have a corresponding argument in `LongformerForTokenClassification.forward` and have been ignored: offset_mapping, overflow_to_sample_mapping.
***** Running Evaluation *****
    Num examples = 1313
    Batch size = 4
Saving model checkpoint to longformer-base-4096-finetuned-ner/checkpoint-1472
Configuration saved in longformer-base-4096-finetuned-ner/checkpoint-1472/config.json
Model weights saved in longformer-base-4096-finetuned-ner/checkpoint-1472/pytorch_model.bin
tokenizer config file saved in longformer-base-4096-finetuned-ner/checkpoint-1472/tokenizer_config.json
Special tokens file saved in longformer-base-4096-finetuned-ner/checkpoint-1472/special_tokens_map.json

```

The following columns in the evaluation set don't have a corresponding argument in `LongformerForTokenClassification.forward` and have been ignored: offset_mapping, overflow_to_sample_mapping.

***** Running Evaluation *****

Num examples = 1313

Batch size = 4

Saving model checkpoint to longformer-base-4096-finetuned-ner/checkpoint-1840

Configuration saved in longformer-base-4096-finetuned-ner/checkpoint-1840/config.json

Model weights saved in longformer-base-4096-finetuned-ner/checkpoint-1840/pytorch_model.bin

tokenizer config file saved in longformer-base-4096-finetuned-ner/checkpoint-1840/tokenizer_config.json

Special tokens file saved in longformer-base-4096-finetuned-ner/checkpoint-1840/special_tokens_map.json

Training completed. Do not forget to share your model on huggingface.co/models =)

Waiting for W&B process to finish... **(success)**.

VBox(children=(Label(value='0.000 MB of 0.000 MB uploaded (0.000 MB deduped)\r'), FloatProgress(value=1.0, max...

Run history:

eval/accuracy	
eval/f1	
eval/loss	
eval/precision	
eval/recall	
eval/runtime	
eval/samples_per_second	
eval/steps_per_second	
train/epoch	
train/global_step	
train/learning_rate	
train/loss	
train/total_flos	
train/train_loss	
train/train_runtime	
train/train_samples_per_second	
train/train_steps_per_second	

Run summary:

eval/accuracy	0.7997
eval/f1	0.22409
eval/loss	0.64129
eval/precision	0.17163
eval/recall	0.32274
eval/runtime	109.5241
eval/samples_per_second	11.988
eval/steps_per_second	3.004
train/epoch	5.0
train/global_step	1840
train/learning_rate	0.0
train/loss	

	0.3923
train/total_flos	3.847061386528358e+16
train/train_loss	0.60354
train/train_runtime	14366.7392
train/train_samples_per_second	4.099
train/train_steps_per_second	0.128

You can sync this run to the cloud by running:

`wandb sync /kaggle/working/wandb/offline-run-20220505_204903-2kdchg1t`

Find logs at: `./wandb/offline-run-20220505_204903-2kdchg1t/logs`

In [30]: `trainer.save_model(model_path)`

Saving model checkpoint to longformer-base-4096-4
 Configuration saved in longformer-base-4096-4/config.json
 Model weights saved in longformer-base-4096-4/pytorch_model.bin
 tokenizer config file saved in longformer-base-4096-4/tokenizer_config.json
 Special tokens file saved in longformer-base-4096-4/special_tokens_map.json

Test Data

```
In [31]: def tokenize_for_test(examples):

    o = tokenizer(examples['text'], truncation=True, return_offsets_mapping=True)

    offset_mapping = o["offset_mapping"]

    o["labels"] = []

    for i in range(len(offset_mapping)):

        labels = [12i['O'] for i in range(len(o['input_ids'][i]))]

        for label_start, label_end, label in \
            list(zip(examples['starts'][i], examples['ends'][i], examples['classifications'][i])):
            for j in range(len(labels)):
                token_start = offset_mapping[i][j][0]
                token_end = offset_mapping[i][j][1]
                if token_start == label_start:
                    labels[j] = 12i[f'B-{label}']
                if token_start > label_start and token_end <= label_end:
                    labels[j] = 12i[f'I-{label}']

        for k, input_id in enumerate(o['input_ids'][i]):
            if input_id in [0,1,2]:
                labels[k] = -100

        labels = fix_beginnings(labels)

        o["labels"].append(labels)

    return o
```



```
In [32]: tokenized_test = datasets.map(tokenize_for_test, batched=True)
tokenized_test
```

```
0%|          | 0/12 [00:00<?, ?ba/s]
0%|          | 0/2 [00:00<?, ?ba/s]
```

```
Out[32]: DatasetDict({
  train: Dataset({
    features: ['__index_level_0__', 'attention_mask', 'classlist', 'ends',
'id', 'input_ids', 'labels', 'offset_mapping', 'predictionstrings', 'starts',
'text'],
    num_rows: 11462
  })
  test: Dataset({
    features: ['__index_level_0__', 'attention_mask', 'classlist', 'ends',
'id', 'input_ids', 'labels', 'offset_mapping', 'predictionstrings', 'starts',
'text'],
    num_rows: 1274
  })
})
```

```
In [33]: # ground truth for test data
```

```
l = []
for example in tokenized_test['test']:
    for c, p in list(zip(example['classlist'], example['predictionstrings'])):
        l.append({
            'id': example['id'],
            'discourse_type': c,
            'predictionstring': p,
        })

gt_df = pd.DataFrame(l)
```

```
In [34]: path = Path('../input/feedback-prize-2021/train')
```

```
colors = {
    'Lead': '#8000ff',
    'Position': '#2b7ff6',
    'Evidence': '#2adddd',
    'Claim': '#80ffb4',
    'Concluding Statement': 'd4dd80',
    'Counterclaim': '#ff8042',
    'Rebuttal': '#ff0000',
    'Other': '#007f00',
}

def visualize(df, text):
    ents = []
    example = df['id'].loc[0]

    for i, row in df.iterrows():
        ents.append({
            'start': int(row['discourse_start']),
            'end': int(row['discourse_end']),
            'label': row['discourse_type']
        })

    doc2 = {
        "text": text,
```

```

        "ents": ents,
        "title": example
    }

    options = {"ents": train.discourse_type.unique().tolist() + ['Other'], "col
displacy.render(doc2, style="ent", options=options, manual=True, jupyter=Tr

```

In [35]: `predictions, labels, _ = trainer.predict(tokenized_test['test'])`

The following columns in the test set don't have a corresponding argument in `LongformerForTokenClassification.forward` and have been ignored: id, ends, start, offset_mapping, __index_level_0__, classlist, predictionstrings, text.
 ***** Running Prediction *****

Num examples = 1274
 Batch size = 4
 Input ids are automatically padded from 462 to 512 to be a multiple of `config.attention_window`: 512

[319/319 01:22]

Input ids are automatically padded from 513 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 944 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1649 to 2048 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 749 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 637 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 935 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 694 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 824 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 469 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 825 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 649 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 786 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 535 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 596 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 586 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 709 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 836 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 575 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 777 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 742 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 935 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 921 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 976 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 623 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 968 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 520 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 632 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 527 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 549 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 629 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1159 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 486 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 526 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 643 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 655 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 535 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1130 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 941 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 546 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 683 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 878 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 677 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 531 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 595 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 834 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 469 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 350 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 638 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 650 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 369 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 781 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 552 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 710 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1104 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 446 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 753 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 805 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 596 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 870 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 690 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 755 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 687 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 434 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 509 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 858 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 518 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 593 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 720 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 430 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 369 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 772 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 820 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 449 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 586 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 697 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 545 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 924 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 654 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 710 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 561 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 897 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1125 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 408 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 641 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 592 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1135 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 662 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 803 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 618 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 772 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 788 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 728 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 758 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 574 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 812 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 788 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 507 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 841 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 549 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 948 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 624 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 582 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1272 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 715 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 500 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 624 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 695 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 656 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 342 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1179 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 595 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 582 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 547 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 992 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 943 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 768 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 768 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 680 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 740 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 621 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 464 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 602 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 563 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 553 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 661 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 555 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 671 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 528 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 589 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 739 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 622 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 651 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 661 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 589 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 810 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 391 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 650 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 574 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 599 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 798 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 988 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 860 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 486 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 469 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 573 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 721 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 918 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 481 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 813 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 423 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 919 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 3991 to 4096 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 775 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 973 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 774 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 692 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1027 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 432 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 745 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 591 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 663 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 777 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 729 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 610 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 419 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 786 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 563 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 807 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 666 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 519 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 859 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 545 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 908 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 570 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 622 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 714 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 733 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1074 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 591 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1242 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 746 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 425 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 703 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 807 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 606 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 854 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 741 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 869 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 437 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 453 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 689 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 653 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 782 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 340 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 516 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1154 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 758 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 597 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 759 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 884 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 566 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 626 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 448 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 745 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 597 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 528 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 814 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 507 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 848 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 749 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1090 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 352 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 859 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 563 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 797 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 390 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 590 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 719 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1027 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 2349 to 2560 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 853 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 654 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 647 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 600 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 501 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 722 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 683 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 629 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 827 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 966 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 485 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 656 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 497 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 565 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1236 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 599 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 784 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 525 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 417 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 895 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 863 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 549 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 727 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 480 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 670 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 618 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 838 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 707 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 568 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 979 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 647 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1068 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 495 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 569 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 422 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 605 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 612 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1044 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1131 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 452 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 577 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 750 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 438 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 840 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 473 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 967 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 386 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 504 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 681 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 787 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 469 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1193 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 755 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 666 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 654 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 670 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 594 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 822 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 711 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 667 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 814 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 484 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 670 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 544 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 475 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 640 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1639 to 2048 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 599 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 863 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 942 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 466 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 519 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 1165 to 1536 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 383 to 512 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 898 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 582 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 770 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 659 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 522 to 1024 to be a multiple of `config.attention_window`: 512

Input ids are automatically padded from 355 to 512 to be a multiple of `config.attention_window`: 512

```

Input ids are automatically padded from 1022 to 1024 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 548 to 1024 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 633 to 1024 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 698 to 1024 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 491 to 512 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 1059 to 1536 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 565 to 1024 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 901 to 1024 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 870 to 1024 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 647 to 1024 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 444 to 512 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 495 to 512 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 536 to 1024 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 381 to 512 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 509 to 512 to be a multiple of `config.attention_window`: 512
Input ids are automatically padded from 467 to 512 to be a multiple of `config.attention_window`: 512

```

```

In [36]: preds = np.argmax(predictions, axis=-1)
         preds.shape

```

```

Out[36]: (1274, 4096)

```

```

In [37]: def get_class(c):
         if c == 14: return 'Other'
         else: return i2l[c][2:]

         def pred2span(pred, example, viz=False, test=False):
             example_id = example['id']
             n_tokens = len(example['input_ids'])
             classes = []
             all_span = []
             for i, c in enumerate(pred.tolist()):
                 if i == n_tokens-1:
                     break
                 if i == 0:
                     cur_span = example['offset_mapping'][i]
                     classes.append(get_class(c))
                 elif i > 0 and (c == pred[i-1] or (c-7) == pred[i-1]):
                     cur_span[1] = example['offset_mapping'][i][1]
                 else:
                     all_span.append(cur_span)
                     cur_span = example['offset_mapping'][i]
                     classes.append(get_class(c))

```

```

all_span.append(cur_span)

if test: text = get_test_text(example_id)
else: text = get_raw_text(example_id)

# abra ka dabra se soli fanta ko pelo

# map token ids to word (whitespace) token ids
predstrings = []
for span in all_span:
    span_start = span[0]
    span_end = span[1]
    before = text[:span_start]
    token_start = len(before.split())
    if len(before) == 0: token_start = 0
    elif before[-1] != ' ': token_start -= 1
    num_tkns = len(text[span_start:span_end+1].split())
    tkns = [str(x) for x in range(token_start, token_start+num_tkns)]
    predstring = ' '.join(tkns)
    predstrings.append(predstring)

rows = []
for c, span, predstring in zip(classes, all_span, predstrings):
    e = {
        'id': example_id,
        'discourse_type': c,
        'predictionstring': predstring,
        'discourse_start': span[0],
        'discourse_end': span[1],
        'discourse': text[span[0]:span[1]+1]
    }
    rows.append(e)

df = pd.DataFrame(rows)
df['length'] = df['discourse'].apply(lambda t: len(t.split()))

# short spans are likely to be false positives, we can choose a min number
df = df[df.length > min_tokens].reset_index(drop=True)
if viz: visualize(df, text)

return df

```

```

In [38]: dfs = []
for i in range(len(tokenized_test['test'])):
    dfs.append(pred2span(preds[i], tokenized_test['test'][i]))

pred_df = pd.concat(dfs, axis=0)
pred_df['class'] = pred_df['discourse_type']
pred_df

```

Out[38]:

	id	discourse_type	predictionstring	discourse_start	discourse_end	discourse
0	F4FD84517F40	Lead	0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18...	1	368	chools are offering distant learning for stude...
1	F4FD84517F40	Position	62 63 64 65 66 67 68 69 70 71 72 73 74 75 76	369	457	Online classes would help tons of students mov...
2	F4FD84517F40	Claim	76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 9...	460	613	Most of the people failing classes or didn't g...
3	F4FD84517F40	Evidence	104 105 106 107 108 109 110 111 112 113 114 11...	614	1367	A big example of this is working. They have to...
4	F4FD84517F40	Claim	237 238 239 240 241 242 243 244 245 246 247 24...	1368	1460	Another case where distant learning may come i...
...
1	58F8F0F77817	Evidence	18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 3...	100	463	"Its a ooc opportunity to take away stress and...
2	58F8F0F77817	Evidence	94 95 96 97 98 99 100 101 102 103 104 105 106 ...	501	1021	you can use someonjes elses gas instead of you...
3	58F8F0F77817	Other	195 196 197 198 199 200 201 202 203 204	1026	1089	iting car usage alsc relieves stress for some ...
4	58F8F0F77817	Evidence	205 206 207 208 209 210 211 212 213 214 215 21...	1090	1464	Some people that i know get really flusterec e...

	id	discourse_type	predictionstring	discourse_start	discourse_end	discourse
			276 277 278 279			limited car
			280 281 282 283			usage is
5	58F8F0F77817	Concluding Statement	284 285 286	1479	1561	gooc
			28...			because it
						save
						sgas...

12284 rows x 8 columns

```
In [39]: def calc_overlap(row):
#Calculates the overlap between prediction and ground truth and
#overlap percentages used for determining true positives.
set_pred = set(row.predictionstring_pred.split(" "))
set_gt = set(row.predictionstring_gt.split(" "))
# Length of each and intersection
len_gt = len(set_gt)
len_pred = len(set_pred)
inter = len(set_gt.intersection(set_pred))
overlap_1 = inter / len_gt
overlap_2 = inter / len_pred
return [overlap_1, overlap_2]

def score_feedback_comp_micro(pred_df, gt_df):
gt_df = (
    gt_df[["id", "discourse_type", "predictionstring"]]
    .reset_index(drop=True)
    .copy()
)
pred_df = pred_df[["id", "class", "predictionstring"]].reset_index(drop=True)
pred_df["pred_id"] = pred_df.index
gt_df["gt_id"] = gt_df.index
# Step 1. all ground truths and predictions for a given class are compared.
joined = pred_df.merge(
    gt_df,
    left_on=["id", "class"],
    right_on=["id", "discourse_type"],
    how="outer",
    suffixes=("_pred", "_gt"),
)
joined["predictionstring_gt"] = joined["predictionstring_gt"].fillna(" ")
joined["predictionstring_pred"] = joined["predictionstring_pred"].fillna(" ")

joined["overlaps"] = joined.apply(calc_overlap, axis=1)

# 2. If the overlap between the ground truth and prediction is >= 0.5,
# and the overlap between the prediction and the ground truth >= 0.5,
# the prediction is a match and considered a true positive.
# If multiple matches exist, the match with the highest pair of overlaps is
joined["overlap1"] = joined["overlaps"].apply(lambda x: eval(str(x))[0])
joined["overlap2"] = joined["overlaps"].apply(lambda x: eval(str(x))[1])

joined["potential_TP"] = (joined["overlap1"] >= 0.5) & (joined["overlap2"] >= 0.5)
joined["max_overlap"] = joined[["overlap1", "overlap2"]].max(axis=1)
tp_pred_ids = (
    joined.query("potential_TP")
    .sort_values("max_overlap", ascending=False)
)
```



```

        .groupby(["id", "predictionstring_gt"])
        .first()["pred_id"]
        .values
    )

    # 3. Any unmatched ground truths are false negatives
    # and any unmatched predictions are false positives.
    fp_pred_ids = [p for p in joined["pred_id"].unique() if p not in tp_pred_ids]

    matched_gt_ids = joined.query("potential_TP")["gt_id"].unique()
    unmatched_gt_ids = [c for c in joined["gt_id"].unique() if c not in matched_gt_ids]

    # Get numbers of each type
    TP = len(tp_pred_ids)
    FP = len(fp_pred_ids)
    FN = len(unmatched_gt_ids)
    # calc microfl
    my_f1_score = TP / (TP + 0.5 * (FP + FN))
    return my_f1_score

def score_feedback_comp(pred_df, gt_df, return_class_scores=False):
    class_scores = {}
    pred_df = pred_df[["id", "class", "predictionstring"]].reset_index(drop=True)
    for discourse_type, gt_subset in gt_df.groupby("discourse_type"):
        pred_subset = (
            pred_df.loc[pred_df["class"] == discourse_type]
            .reset_index(drop=True)
            .copy()
        )
        class_score = score_feedback_comp_micro(pred_subset, gt_subset)
        class_scores[discourse_type] = class_score
    f1 = np.mean([v for v in class_scores.values()])
    if return_class_scores:
        return f1, class_scores
    return f1

```

F-1 Score on Test Data

```
In [40]: score_feedback_comp(pred_df, gt_df, return_class_scores=True)
```

```
Out[40]: (0.6006660382605588,
          {'Claim': 0.5415699024616814,
           'Concluding Statement': 0.7681779298545766,
           'Counterclaim': 0.4782608695652174,
           'Evidence': 0.6717257621352872,
           'Lead': 0.7335858585858586,
           'Position': 0.6224530563324011,
           'Rebuttal': 0.3888888888888889})
```

Demo - Visualising Predictions on Test Data

```
In [44]: pred_df.head()
```

Out [44]:

	id	discourse_type	predictionstring	discourse_start	discourse_end	discourse
0	F4FD84517F40	Lead	0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18...	1	368	chools are offering distant learning for stude...
1	F4FD84517F40	Position	62 63 64 65 66 67 68 69 70 71 72 73 74 75 76	369	457	Online classes would help tons of students mov...
2	F4FD84517F40	Claim	76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 9...	460	613	Most of the people failing classes or didn't g...
3	F4FD84517F40	Evidence	104 105 106 107 108 109 110 111 112 113 114 11...	614	1367	A big example of this is working. They have to...
4	F4FD84517F40	Claim	237 238 239 240 241 242 243 244 245 246 247 24...	1368	1460	Another case where distant learning may come i...

In [42]:

pred2span(preds[0], tokenized_test['test'][0], viz=True)

F4FD84517F40

Schools are offering distant learning for students to attend classes off campus. Many students are saying this is great and they are taking the offer to continue and better their education. On the other hand, some say that this is not the best idea for schooling. Schools want to give students an opportunity to better their learning and stay connected with education. **Lead** Online classes would help tons of students move closer to graduation and a better future **Position** .

Most of the people failing classes or didn't graduate the first time around have things they have to do outside of school to help support their families. **Claim** A big example of this is working. They have to work to make money for their families but without all of the education they need, they can't work for a better paying job. Therefore, lots of people are working multiple jobs to support families; this is leaving them no time to spend six to eight hours of the day in a building. Another example is that nowadays many teens are having to stay home to watch their own child or watching a family member because of finances in their home. Although being home may help out with home situations, students grades are dropping due to the amount of absent days in classes. Online classes and video conferencing would benefit these students by allowing them to take classes at night and/or multitask during the day. **Evidence**

Another case where distant learning may come in use is for students very involved in sports. **Claim** Doing a sport can be a big responsibility. Sports are very time consuming, although some athletes make it look easy, there is a lot that goes into playing. Many levels of sports may require leaving the state or country for days, sometimes even weeks. While playing a sport it is easy to be injured or become sick. Furthermore, students are losing class time and unable to make up for missed days in little time. Distant learning would be beneficial for those who travel so they can take their school work along and manage their free time. Online classes also help when they are sick or injured to give them time to rest and heal and still not miss out on continuing to learning. **Evidence**

In conclusion, yes distant learning as an option to students is very beneficial. Allowing students to attend class from home or while traveling could get more and more people willing to go further in learning. **Concluding Statement**

Out [42]:

	id	discourse_type	predictionstring	discourse_start	discourse_end	discourse
0	F4FD84517F40	Lead	0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18...	1	368	chools a offerir dista learning f stude
1	F4FD84517F40	Position	62 63 64 65 66 67 68 69 70 71 72 73 74 75 76	369	457	Onlir classes wou help tons studen mov
2	F4FD84517F40	Claim	76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 9...	460	613	Most of th people failir classes didn't g
3	F4FD84517F40	Evidence	104 105 106 107 108 109 110 111 112 113 114 11...	614	1367	A b example this workin They ha to
4	F4FD84517F40	Claim	237 238 239 240 241 242 243 244 245 246 247 24...	1368	1460	Another ca: where dista learning m: come i
5	F4FD84517F40	Evidence	254 255 256 257 258 259 260 261 262 263 264 26...	1461	2141	Doing a spc can be a b responsibilit Spo
6	F4FD84517F40	Concluding Statement	376 377 378 379 380 381 382 383 384 385 386 38...	2142	2351	In conclusio yes dista learning as : opti

In []:

In []: